

Call Intelligence Summarizer

Transcribes discovery calls, scores them against a qualification framework, and writes structured notes and next steps back to the CRM.

Python	Whisper	Claude API	HubSpot
--------	---------	------------	---------

VERSION	STATUS	STAGE	RUNTIME	OWNER	UPDATED
0.6.0	In Progress	Pilot, 2 teams	Python 3.12	Sales Eng	Sep 2026

1. Overview

The Call Intelligence Summarizer takes a recorded discovery call and returns a structured record of it: a transcript with speakers separated, a qualification score against the team's framework, the notes a rep should have written, and the next steps that were actually agreed. Those artifacts are written back to HubSpot against the right contact and deal, so the CRM reflects the conversation without the rep retyping it.

Scoring is the part that earns its place. A discovery call either surfaced the metric, the economic buyer, the decision process, and the compelling event, or it did not; the model marks each dimension as covered, partially covered, or missed, and cites the transcript span that justifies the mark. Managers coach against the gaps instead of against a recollection, and a rep can contest a score by reading the cited line.

2. Processing path




Figure 1 — Redaction sits before analysis, so raw payment and identity data never reaches the model provider.

Every stage is checkpointed to object storage under the call ID, so a failure in analysis does not force a re-transcription — the expensive step — and prompt changes can be replayed over historical transcripts to compare scoring behaviour before a new version is promoted.

3. Components

Ingest	Webhook from the dialer and meeting recorder; pulls audio, deduplicates re-sends by call ID.
Transcription	Whisper large-v3 with speaker diarization; word-level timestamps retained for citation.
Redaction	Rule and NER-based scrub of card numbers, national IDs, and personal contact details.
Analysis	Claude API with a JSON schema for summary, qualification marks, objections, and next steps.
Write-back	HubSpot engagement, structured deal properties, and tasks for each owned next step.
Review queue	Low-confidence or contested outputs held for human confirmation before write-back.

4. Output schema

```
{
  "call_id": "c_88213", "deal_id": "14420031", "duration_s": 2714,
  "summary": "Evaluating a replacement for an in-house tool ahead of ...",
  "qualification": {
    "metric": { "state": "covered", "span": [412, 468] },
    "economic_buyer": { "state": "partial", "span": [1180, 1204] },
    "decision_process": { "state": "missed", "span": null },
    "compelling_event": { "state": "covered", "span": [2260, 2298] }
  },
  "objections": [{ "text": "Concern about migration effort", "span": [1902, 1944] }],
  "next_steps": [
    { "owner": "rep", "action": "Send security questionnaire", "due": "2026-09-15" },
    { "owner": "prospect", "action": "Confirm procurement contact", "due": "2026-09-18" }
  ],
  "confidence": 0.81
}
```

Listing 1 — Spans are millisecond offsets into the transcript, so every claim links to audio.

5. Current status

CAPABILITY	STATE	NOTES
Transcription and diarization	Shipped	Stable across the dialer and meeting recorder.
Summary and next steps	Shipped	In pilot use by two teams since July.
Qualification scoring	In progress	Agreement with human raters at 0.78; target is 0.85.
Multi-language support	Not started	English only; Spanish scheduled next.
Bidirectional CRM sync	Not started	Write-back only; CRM edits do not flow back.

6. Evaluation approach

- **Ground truth.** A rolling set of 200 calls scored independently by two sales managers; disagreements between the humans are excluded rather than averaged, so the benchmark measures only unambiguous cases.
- **Regression gate.** Prompt or model changes must hold or improve per-dimension agreement on that set before promotion, and every version's scores are retained for comparison.
- **Known weakness.** Agreement drops on calls with heavy crosstalk or more than four participants, where diarization errors propagate into attribution of who said what.

7. Privacy and consent

- Recording consent is verified upstream by the dialer; calls without a consent flag are rejected at ingest and never transcribed.
- Redaction runs before any external API call, and audio is never sent to the model provider — only redacted text. Raw audio expires from object storage after 90 days.
- Transcript access follows CRM deal permissions, so a rep cannot read a call on a deal they cannot already see.

8. Remaining work

MILESTONE	SCOPE	TARGET
Pilot exit	Scoring agreement at 0.85, review queue under 10% of calls.	Q4 2026
General release	All sales teams, self-serve framework configuration.	Q1 2027
1.0	Spanish and German, coaching digests, objection trend reporting.	Q2 2027